

Optimizing Performance at UConn

Testing, Analyzing, and Improving KFS Performance

Leo Przybylski ¹ James Gedarovich ³

rSmart¹

leo@rsmart.com

UConn ³

james.gedarovich@uconn.edu

October 17, 2012

About Presenters

- ▶ Jim Gedarovich
- ▶ Leo Przybylski

Overview

- ▶ Planning
 - ▶ Test plan and schedule
 - ▶ Notifying interested parties (get people involved)
- ▶ Infrastructure
 - ▶ Test Exercisers
 - ▶ Application Hosts
- ▶ Tools
 - ▶ Tsung
 - ▶ Recorder
 - ▶ Open Performance Automation Framework
 - ▶ Open-Synchronized-System-Resource-Monitoring
 - ▶ jstat
 - ▶ aragozin/jvm-tools (JTop GCRep)
 - ▶ JVisualVM and AppDynamics

More Overview

- ▶ Building Tests with the Open Performance Automation Framework
- ▶ Running Tests
- ▶ Reporting Statistics
- ▶ Analyzing Test Results
- ▶ Translating Analysis Data into Optimizations
- ▶ Lessons Learned

Tools - Tsung

- ▶ Setup and Installation
- ▶ How it works
 - ▶ Request simulation (Not browser simulation)
 - ▶ Tests the application, not the server
- ▶ Features

Tools - Open Performance Automation Framework

- ▶ DSL for Tsung
- ▶ Makes writing tests very easy
- ▶ Libraries for canned test function
- ▶ Allows scripting of Tsung which creates

Tools - Open-Synchronized-System-Resource-Monitoring

- ▶ Aggregates resource reports (sar, jstat, netstat, etc...) from multiple sources.
- ▶ Utilizes perl modules.
- ▶ Needed to create a module for jstat.
- ▶ Eliminates human error by automating report generation and gathering.

Tools - jstat and jstatd

It's a command line tool

```
r351574nc3@behemoth~
```

```
(13:41:22) [24] jstat
```

```
invalid argument count
```

```
Usage: jstat -help|-options
```

```
    jstat -<option> [-t] [-h<lines>] <vmid> [<interval> [<count>]]
```

Definitions:

<option> An option reported by the -options option
<vmid> Virtual Machine Identifier. A vmid takes the following form:

```
    <lvmid>[@<hostname>[:<port>]]
```

Where <lvmid> is the **local** vm identifier **for** the target Java virtual machine, typically a process id; <hostname> is the name of the host running the target Java virtual machine; and <port> is the port number **for** the rmiregistry on the target host. See the jvmstat documentation **for** a more **complete** description of the Virtual Machine Identifier.

<lines> Number of samples between header lines.

<interval> Sampling interval. The following forms are allowed:

```
    <n>["ms"|"s"]
```

Where <n> is an integer and the suffix specifies the units as milliseconds("ms") or seconds("s"). The default units are "ms".

<count> Number of samples to take before terminating.

-J<flag> Pass <flag> directly to the runtime system.

Tools - jstat and jstatd

Kicking it off

```
r351574nc3@behemoth~
```

```
(13:50:33) [36] jstat 99672 10 720
```

S0	S1	E	O	P	YGC	YGCT	FGC	FGCT	GCT
0.00	0.00	73.14	32.32	53.38	2758	93.932	30	51.206	145.137
0.00	0.00	79.10	32.32	53.38	2758	93.932	30	51.206	145.137
0.00	0.00	86.86	32.32	53.38	2758	93.932	30	51.206	145.137
0.00	0.00	96.22	32.32	53.38	2758	93.932	30	51.206	145.137
0.00	0.00	1.93	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	16.57	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	24.34	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	33.70	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	45.62	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	68.19	33.27	53.38	2759	94.013	30	51.206	145.218
0.00	0.00	79.84	33.27	53.38	2759	94.013	30	51.206	145.218

Tools - aragozin/jvm-tools (JTop GCRep)

- ▶ <https://github.com/aragozin/jvm-tools>
- ▶ Very useful for dumping cpu and gc usage per thread.
- ▶ Much like jstat.
- ▶ More of the information you want all in one place.
- ▶ JTopStats
<https://github.com/ybart/JTopStats> displays information in a web application

Building Tests with the Open Performance Automation Framework

- ▶ The framework
- ▶ <http://www.github.com/leo-at-rsmart/Open-Performance-Automation-Framework>
- ▶ Handling posts

Open Performance Automation Framework Structure

```
|-  
|-/config  
|---data  
|-----kfs  
|-----data  
|---import  
|-----kfs  
|---tests  
|---lib  
|-----kfs  
|---tests  
|-----kfs  
|-/lib  
|---kfs  
|-----common  
|-----dv  
|-/log  
|-/suites  
|---kfs  
|-/tests  
|---kfs
```

Starting with the Tsung Recorder

- ▶ Eliminates human error.
- ▶ Useful for gathering data for multipart form submissions.
- ▶ It's a proxy
- ▶ to start it:

```
(11:24:33) [1] /opt/local/bin/tsung-recorder --help
/opt/local/bin/tsung-recorder: illegal option -- -
Usage: tsung-recorder <options> start|stop|restart
Options:
```

```
-p <plugin>      plugin used for the recorder
                  available: http, pgsql,webdav (default is http)
-L <port>        listening port for the recorder (default is 8090)
-I <IP>          for the pgsql recorder (or parent proxy): server IP
                  (default is 127.0.0.1)
-P <port>        for the pgsql recorder (or parent proxy): server port
                  (default is 5432)
-u              for the http recorder: use a parent proxy
-d <level>       set log level from 0 to 7 (default is 5)
-v              print version information and exit
-h              display this help and exit
```

Starting with the Tsung Recorder

To start

```
tsung-recorder start
```

Don't forget to set the proxy in your browser.

Recorder Results

Stores in $\$HOME/.tsung/$ formatted as
tsung_recorder20120502 – 1006 – 1.bin

-----19277021961952509530130060903
Content-Disposition: form-data; name="tabStates(DocumentOverview)"

OPEN

-----19277021961952509530130060903
Content-Disposition: form-data; name="document.documentHeader.documentNumber"

%%_document_number%%

-----19277021961952509530130060903
Content-Disposition: form-data; name="document.documentHeader.documentDescripti

Duplicating

-----19277021961952509530130060903
Content-Disposition: form-data; name="document.documentHeader.explanation"

Test Development Patterns

- ▶ Develop tests per document
- ▶ Common actions can be moved into libraries and reused
- ▶ Simulate user clicks by adding pauses

Infrastructure/Creating a Test Environment: Exercisors

- ▶ Plan on Unix open file limitations for socket connections
- ▶ Exercisors can consume CPU and hard disk resources
- ▶ Shoot for a minimum of
 - ▶ 2.5Ghz 4 Cores with 4 Gb of memory.
 - ▶ 250 Gb of hard disk space (logs can take up a lot of space).
 - ▶ 2 - 4 exercisors (assume 1 exercisor is equivalent to 50 users).
- ▶ Virtual machines are acceptable.

Infrastructure/Creating a Test Environment: Application Servers

- ▶ Isolate environments.
- ▶ Configure environments with and without balancing.

Caveats

- ▶ Impacting External Live Systems
 - ▶ The boundary between testing environment and user environment
 - ▶ When is it too realistic?
 - ▶ External live systems impact your application and can create latency that needs to be tested
 - ▶ CAS
 - ▶ Kerberoas
 - ▶ LDAP

Testing

- ▶ Testing during peak hours

Analysis

- ▶ Understanding requests vs. connections
- ▶ Data Correlations
- ▶ Results across platform to determine the best hardware

Optimization

- ▶ Improving Request Times
 - ▶ SessionDocumentService
 - ▶ Improved Logging
 - ▶ Reducing data transfered
 - ▶ Reducing web service calls
 - ▶ Transactional vs. Non-transactional datasources
- ▶ Optimal hardware configurations
 - ▶ CPU impact
 - ▶ Memory impact

Lessons Learned

1. Test frequently.
2. Establish a baseline for analysis.
3. Isolate environments for load testing.
4. Determine the impact on users through external systems.
5. Keep support staff informed and available during testing.
6. They are tests. Automate as much as possible to avoid human error.
7. Establish a feedback system to interested parties to resolve issues quickly.

Thank you

We hope this session was informative